

Ganesh Karayi

916-284-4014 | ganeshkarayi@berkeley.edu | [linkedin.com/in/ganesh-karayi](https://www.linkedin.com/in/ganesh-karayi)

EDUCATION

University of California, Berkeley

Expected May 2028

B.A. Applied Mathematics and Data Science, GPA: 3.8

Berkeley, CA

- **Relevant Coursework:** Data Structures & Algorithms, Probability & Random Processes, Principles & Techniques of Data Science, Real Analysis, Intro to Artificial Intelligence (CS 188); **In Progress:** Intro to Machine Learning (CS 189), Optimization Models in Engineering (EECS 127)

EXPERIENCE

IPick.AI

Jan. 2026 – Present

AI Engineer

Cupertino, CA

- Led a team of 14 to ship a live equity intelligence platform covering 4,000+ public companies, visualizing supplier, competitor, and subsidiary relationships as an interactive knowledge graph
- Engineered an NLP extraction pipeline processing 8,000+ SEC EDGAR filings, isolating relevant sections as context and classifying corporate relationships via FinBERT, Claude API, and Llama 3.2
- Directed architecture decisions across the full stack (Python/Flask, PostgreSQL, Docker, AWS EC2, D3.js), scraping 20+ financial news sources and feeding extracted sections into Claude to generate expert-level AI stock analyses surfaced per ticker in real time

Shepherd Ventures

Jan. 2026 – Present

Senior Quantitative Analyst

San Diego, CA

- Built a black swan fear-regime detection system with the Claude API, combining moving-average slope decay, fast-shock drawdown triggers, and exogenous override signals, reaching 100% recall across 8 labeled crash events in a 20-year backtest
- Benchmarked an XGBoost crash-classification baseline (82.9% accuracy, 85.3% stress recall, 30.2% stress precision), diagnosing precision/recall trade-offs via a false-positive/false-negative error ledger spanning 10 historical stress events including COVID and the 2018 Volmageddon selloff

DealPrint

Aug. 2025 – Jan. 2026

Data Science Research Intern

Berkeley, CA

- Independently built an end-to-end M&A valuation tool from data acquisition through modeling to delivery, self-sourcing 50,000 transactions spanning 40 years via a custom-built Chrome extension scraping a legacy web platform
- Engineered high-dimensional features across a tiered KNN comparables engine with adjustable similarity weighting spanning a 2/4/6-digit NAICS industry hierarchy, geography, and deal type, feeding a regression model on revenue, SDE/EBITDA margin, state, and industry features to produce sale price estimates at a MAPE of 0.25 on held-out transactions

PROJECTS

Solar Panel Market Quantile Analysis

2026

- Self-sourced and cleaned a 332,000-row, 266-column California solar interconnection dataset, engineering non-linear size/storage tier features and a cost-per-Watt target variable
- Built a custom quantile regression estimator from scratch, tuning kernel bandwidth for the model's density-weighting step and validating predictive accuracy against statsmodels.QuantReg, surfacing nonlinear cost dynamics including a 90th-percentile price premium and battery-storage economies of scale

ACTIVITIES

Open Project — Vice President of Projects

Aug. 2025 – Present

- Oversee 15+ simultaneous client projects and 40+ Project Managers across the full product lifecycle, from cold outreach and client sourcing through scoping, build, and delivery
- Interviewed and selected 50+ PM candidates, evaluating leadership, communication, and technical fit

SKILLS AND INTERESTS

Languages: Python, SQL, JavaScript, Java

ML & AI: scikit-learn, XGBoost, FinBERT, Llama 3.2, Claude API, NLP, Feature Engineering

Developer Tools & Infrastructure: Git, Docker, AWS (EC2), PostgreSQL, FastAPI, Flask

Interests: Prediction Markets (Kalshi), Poker, Rock Climbing